

# Embedded Systems Design: A Unified Hardware/Software Introduction

---

## Chapter 3 General-Purpose Processors: Software



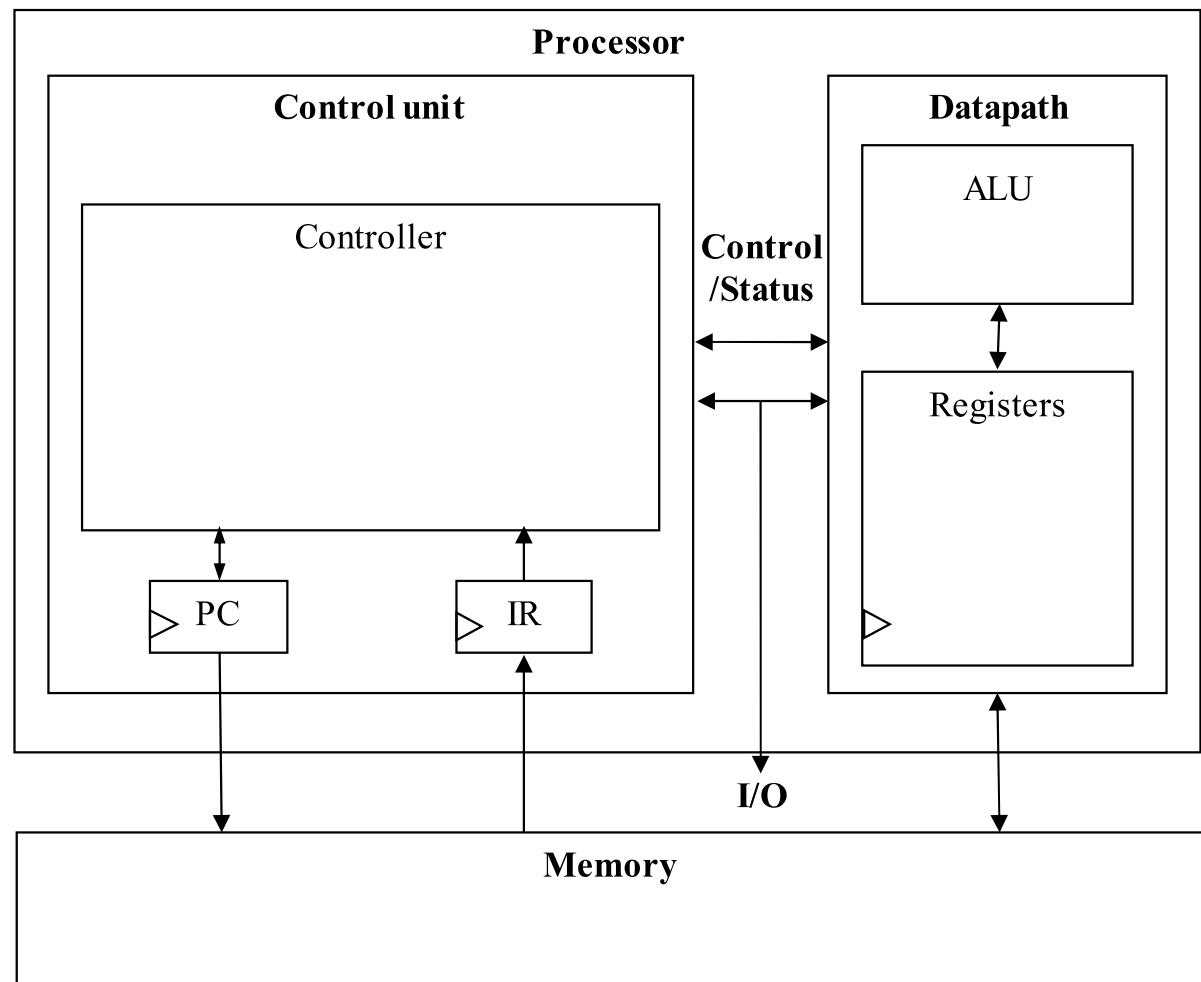
# Introduction

---

- General-Purpose Processor
  - Processor designed for a variety of computation tasks
  - Low unit cost, in part because manufacturer spreads NRE over large numbers of units
    - Motorola sold half a billion 68HC05 microcontrollers *in 1996 alone*
  - Carefully designed since higher NRE is acceptable
    - Can yield good performance, size and power
  - Low NRE cost, short time-to-market/prototype, high flexibility
    - User just writes software; no processor design
  - a.k.a. “microprocessor” – “micro” used when they were implemented on one or a few chips rather than entire rooms

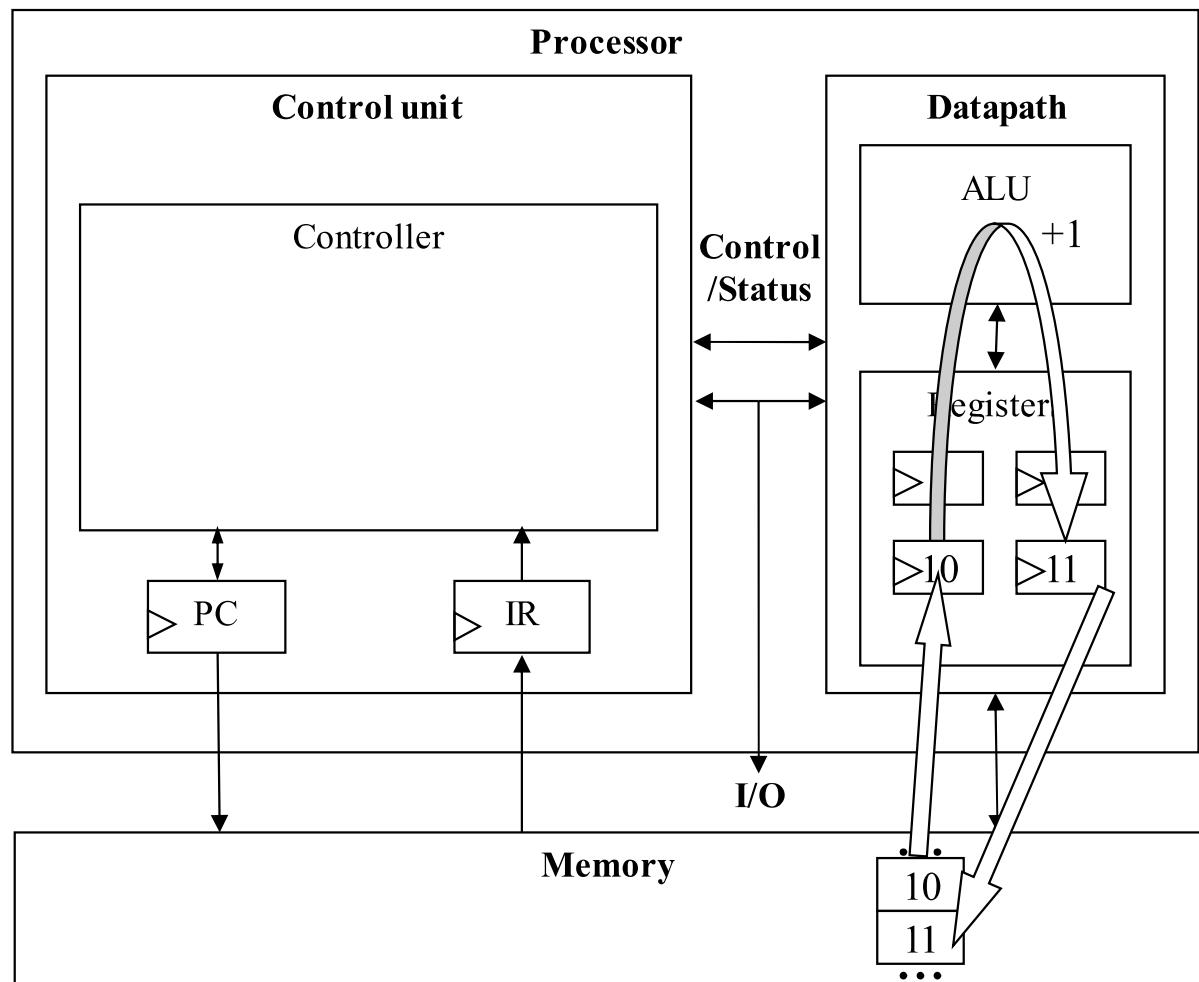
# Basic Architecture

- Control unit and datapath
  - Note similarity to single-purpose processor
- Key differences
  - Datapath is general
  - Control unit doesn't store the algorithm – the algorithm is “programmed” into the memory



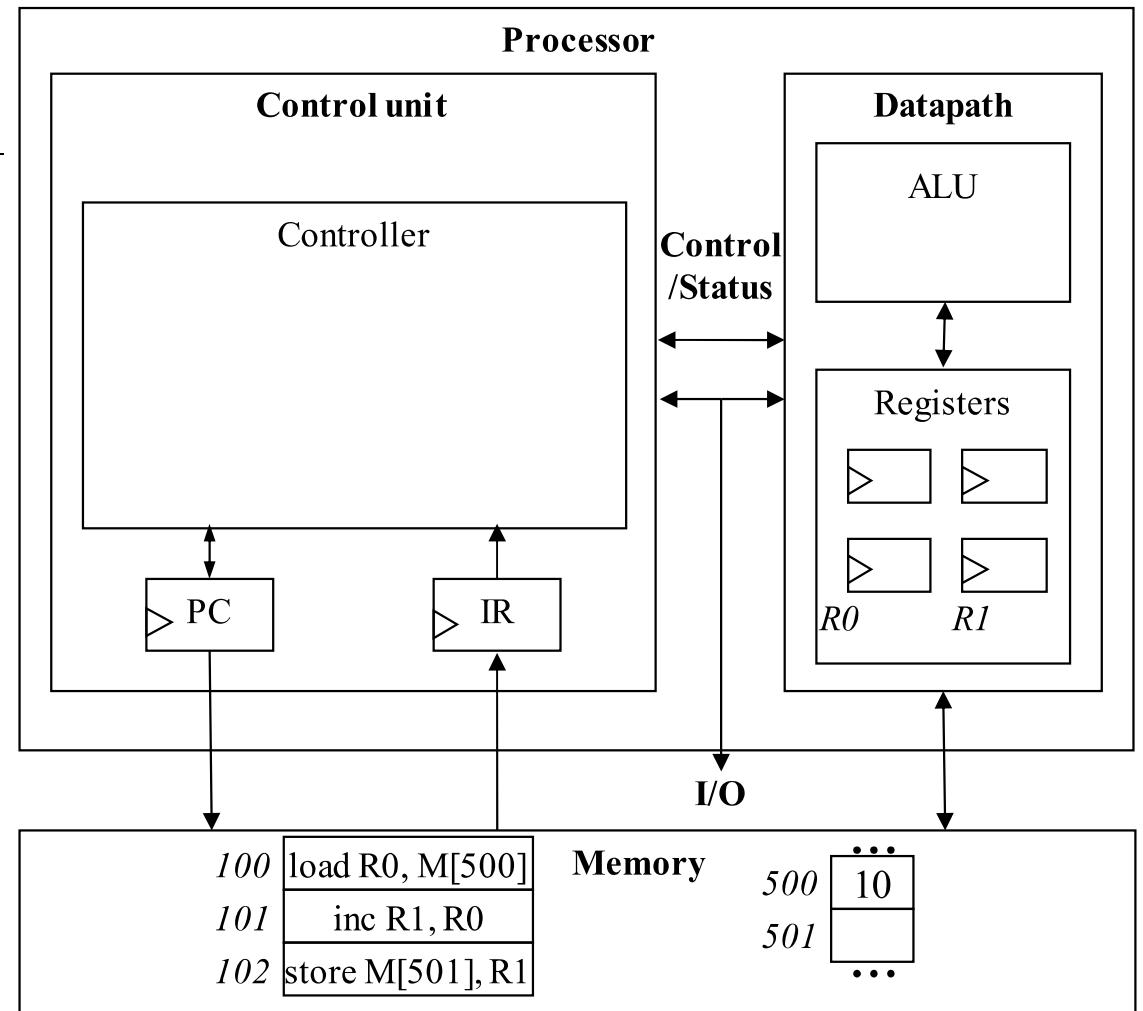
# Datapath Operations

- Load
  - Read memory location into register
- ALU operation
  - Input certain registers through ALU, store back in register
- Store
  - Write register to memory location



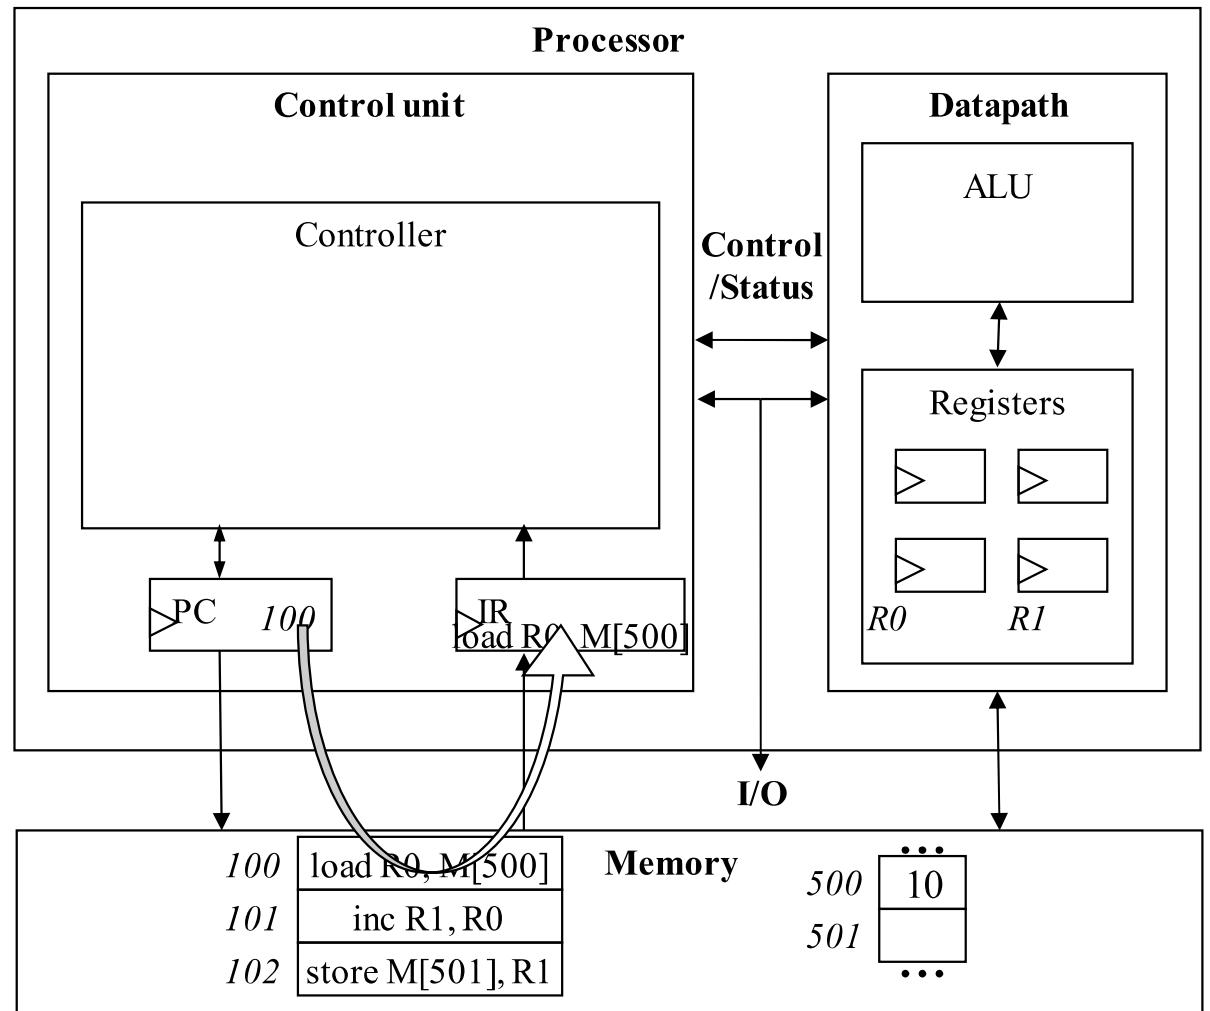
# Control Unit

- Control unit: configures the datapath operations
  - Sequence of desired operations (“instructions”) stored in memory – “program”
- Instruction cycle – broken into several sub-operations, each one clock cycle, e.g.:
  - Fetch: Get next instruction into IR
  - Decode: Determine what the instruction means
  - Fetch operands: Move data from memory to datapath register
  - Execute: Move data through the ALU
  - Store results: Write data from register to memory



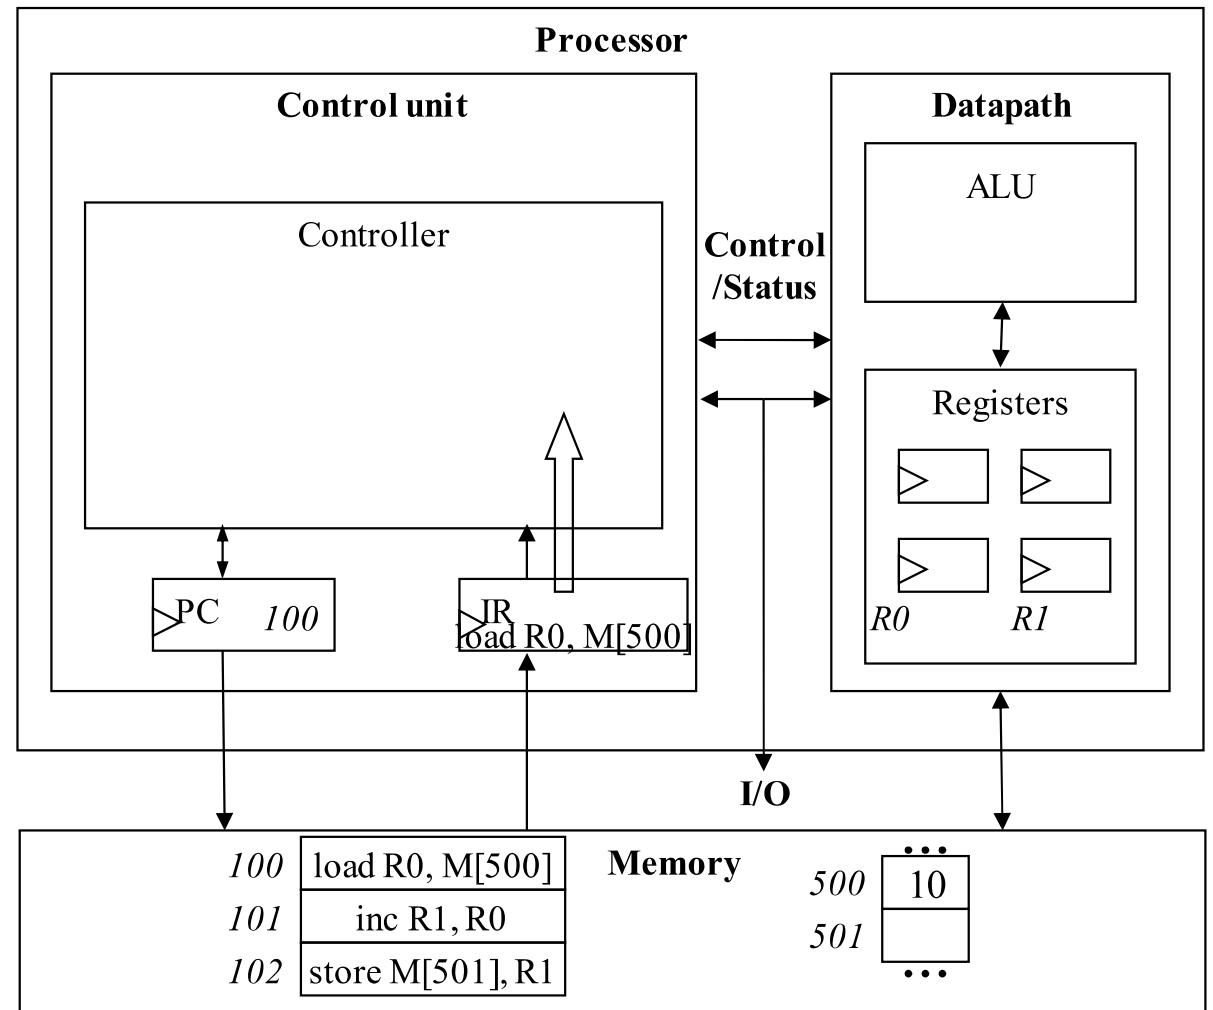
# Control Unit Sub-Operations

- Fetch
  - Get next instruction into IR
  - PC: program counter, always points to next instruction
  - IR: holds the fetched instruction



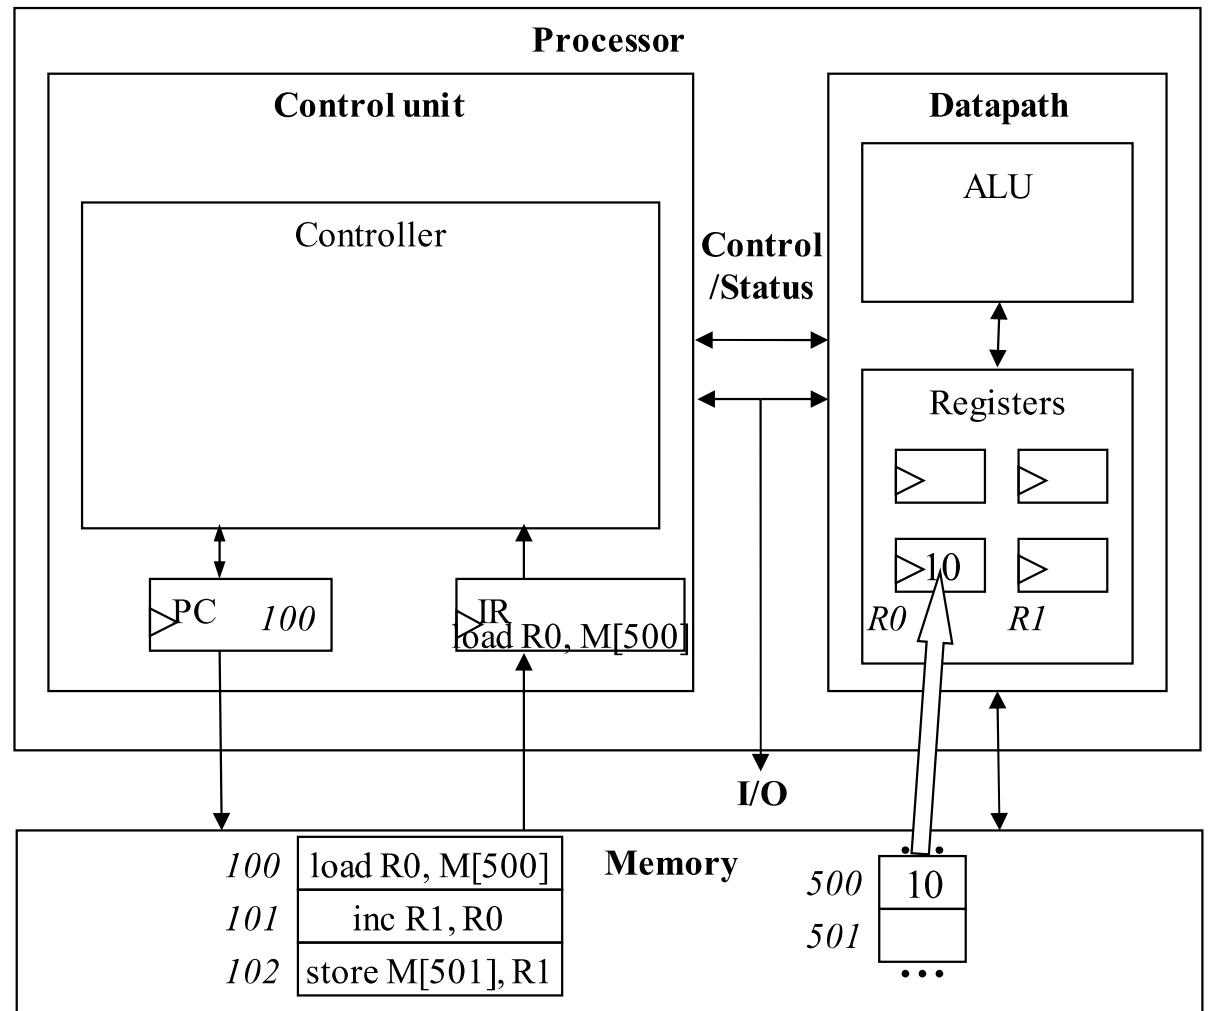
# Control Unit Sub-Operations

- Decode
  - Determine what the instruction means



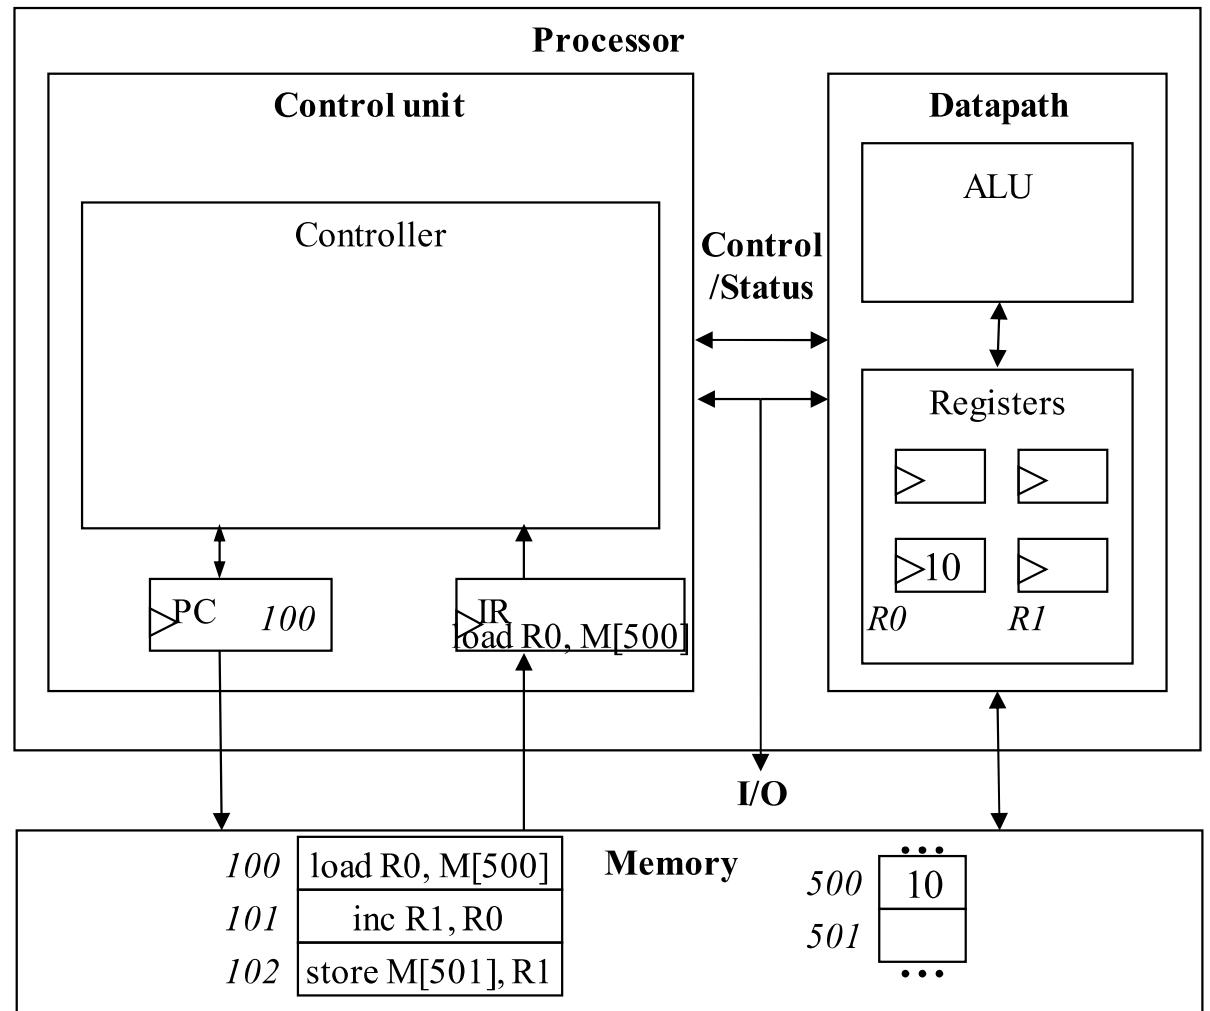
# Control Unit Sub-Operations

- Fetch operands
  - Move data from memory to datapath register



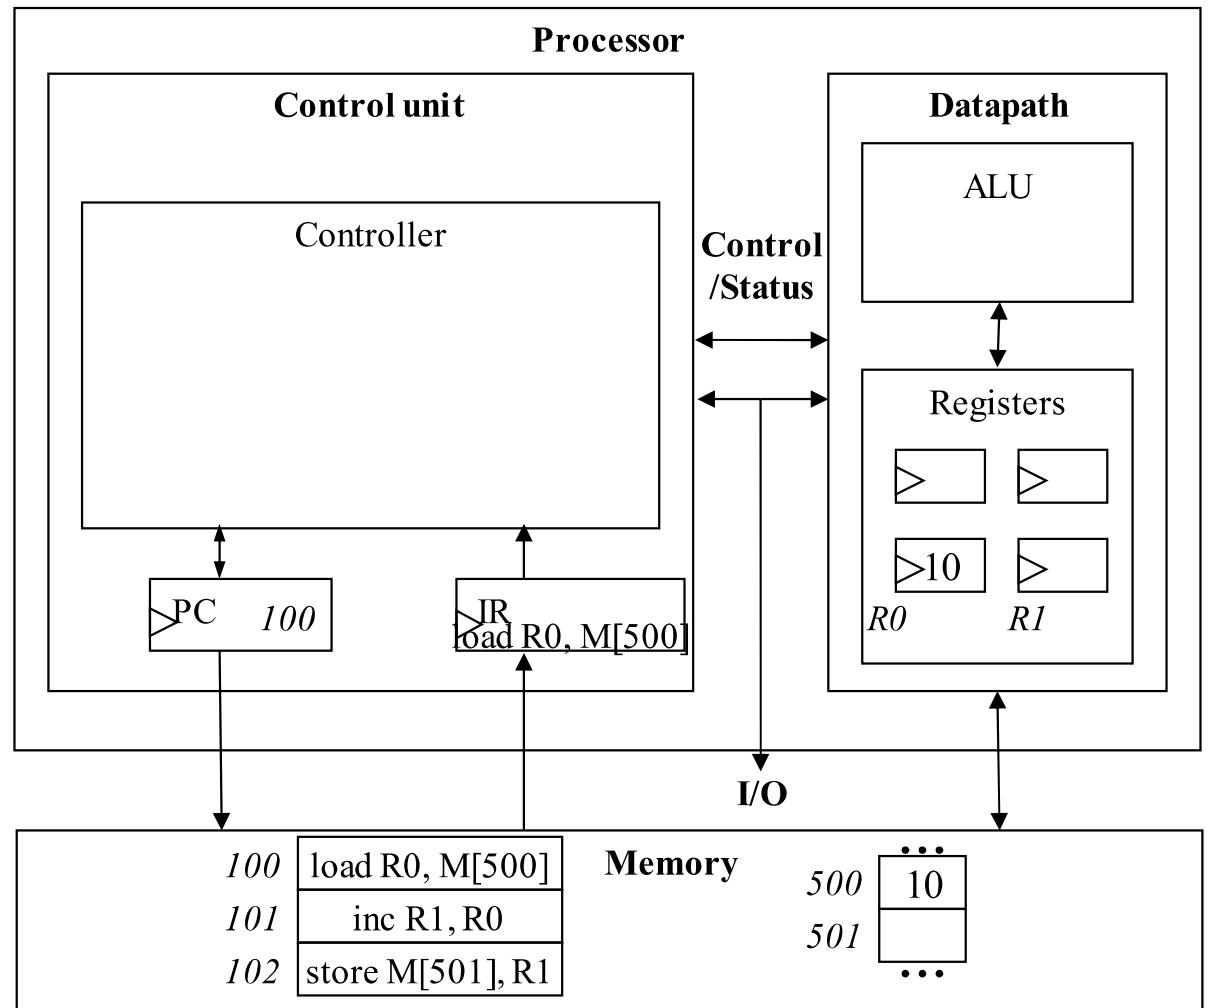
# Control Unit Sub-Operations

- Execute
  - Move data through the ALU
  - This particular instruction does nothing during this sub-operation

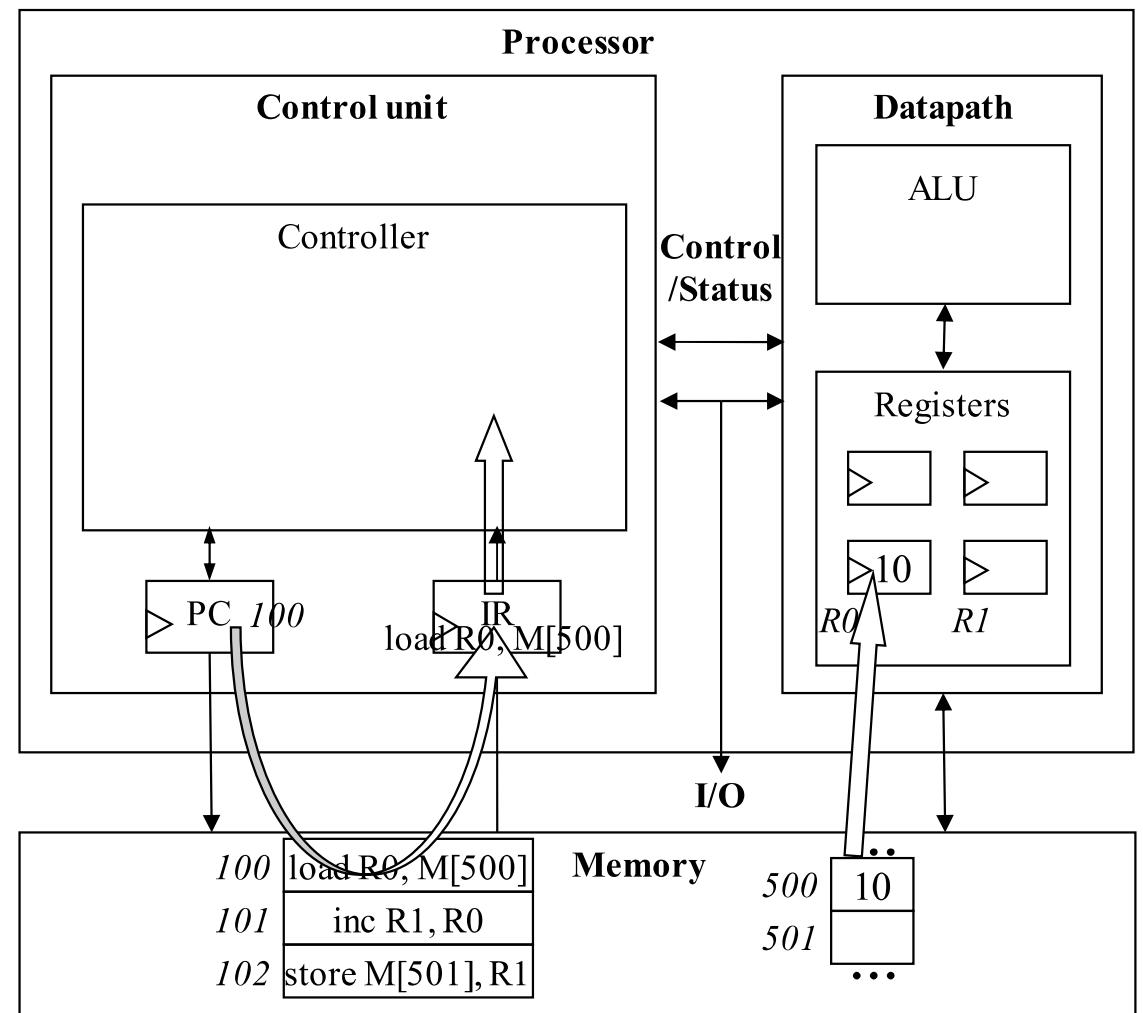
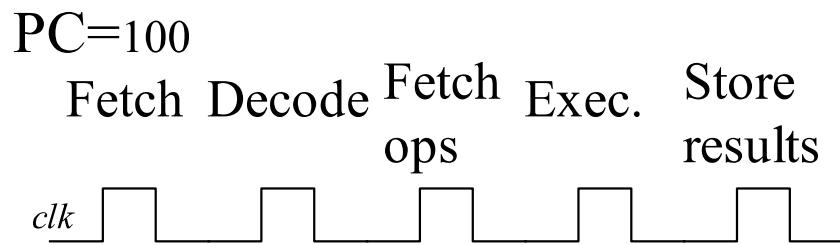


# Control Unit Sub-Operations

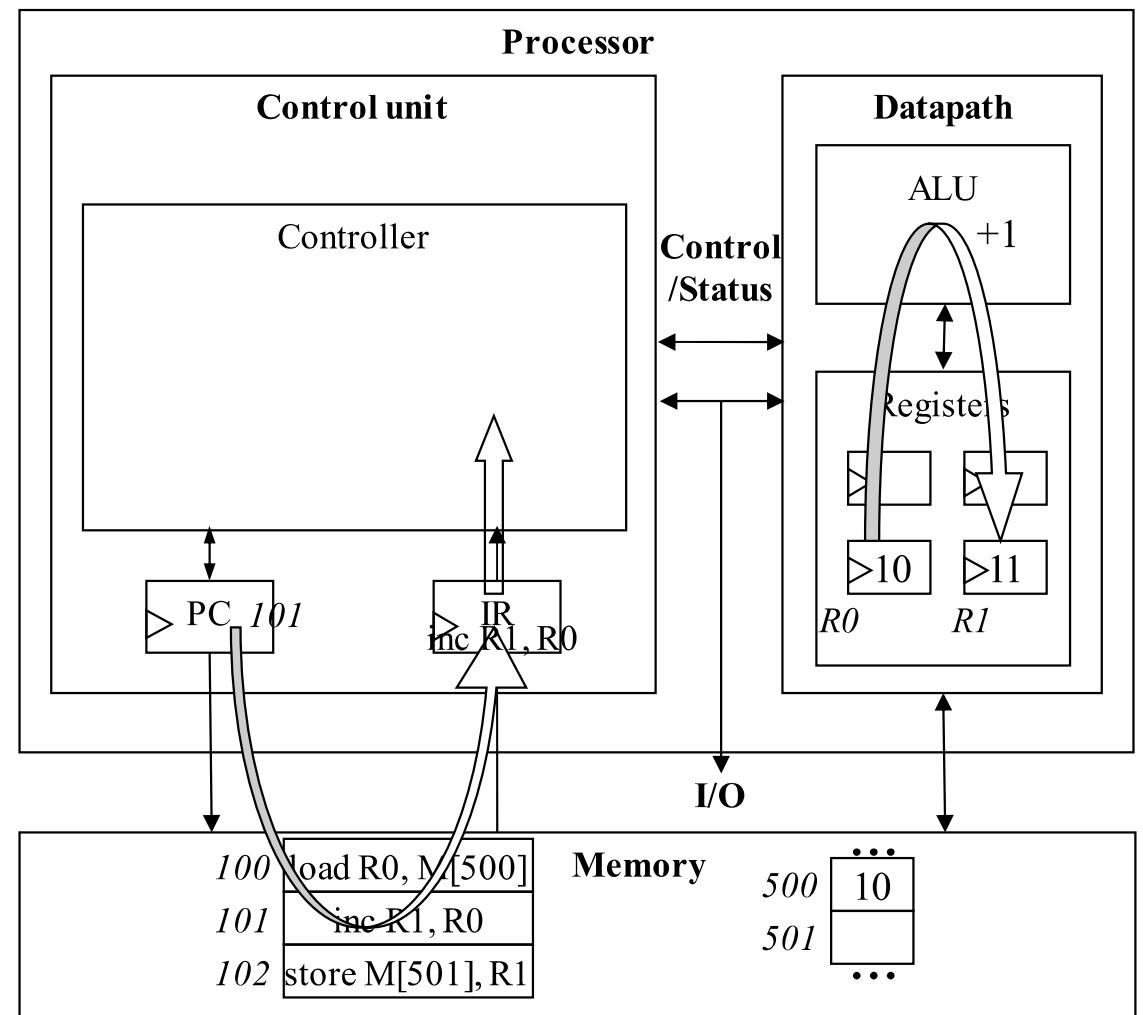
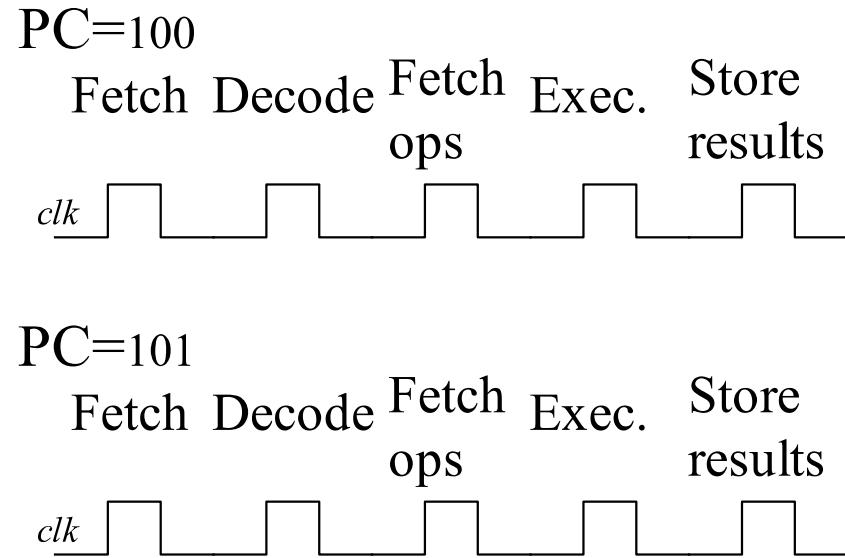
- Store results
  - Write data from register to memory
  - This particular instruction does nothing during this sub-operation



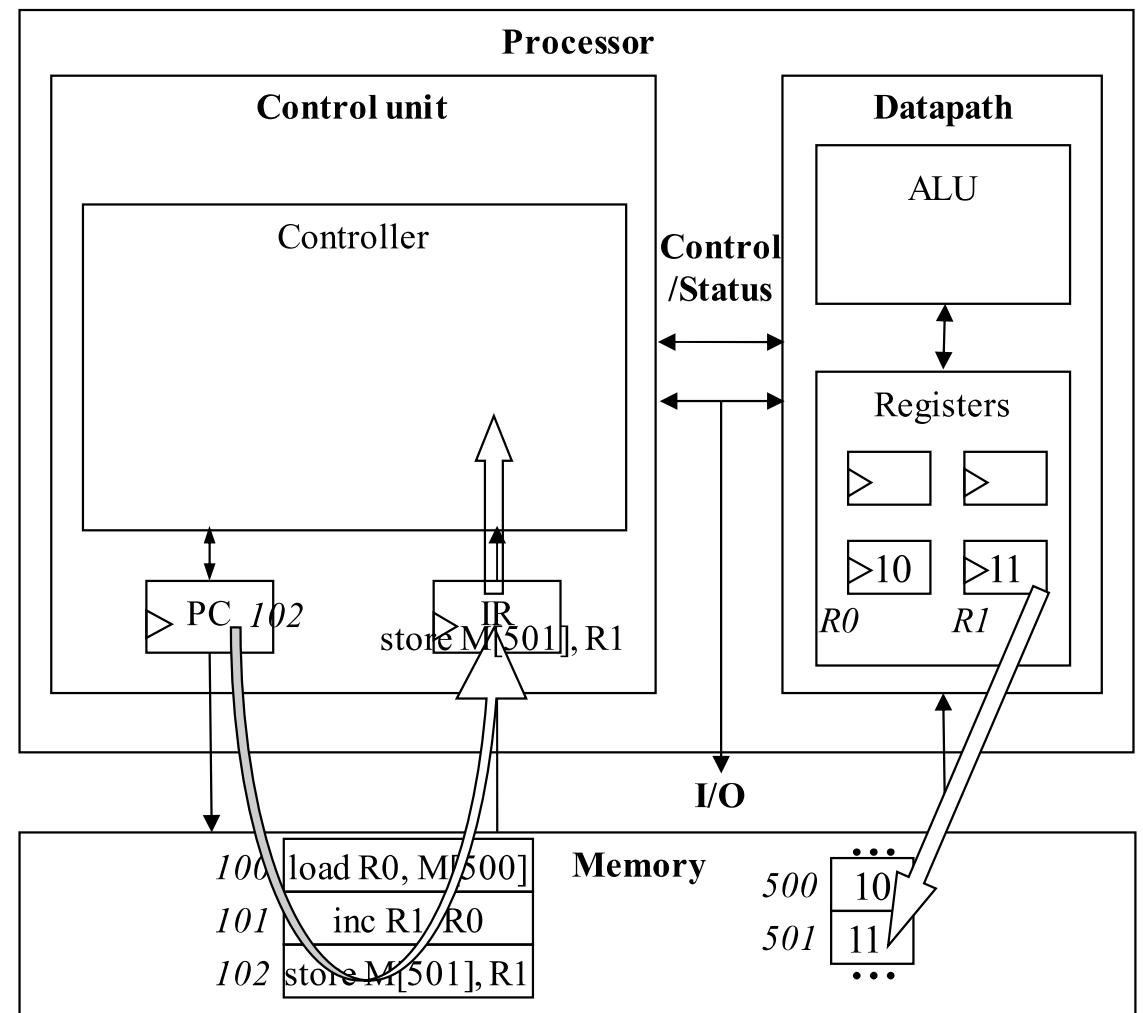
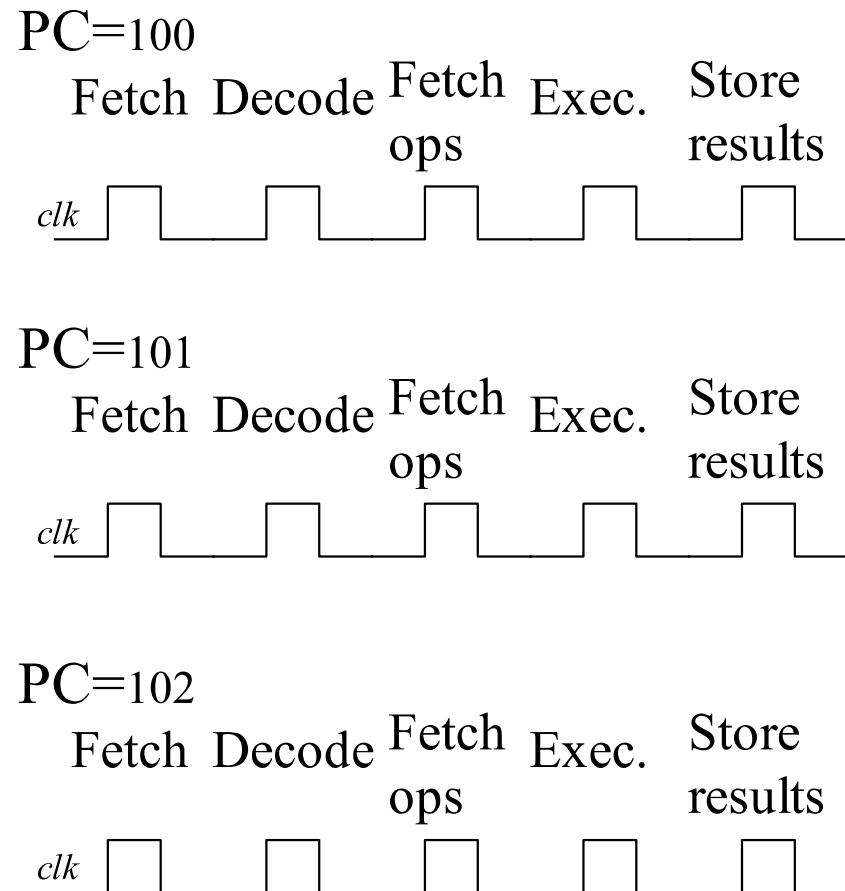
# Instruction Cycles



# Instruction Cycles

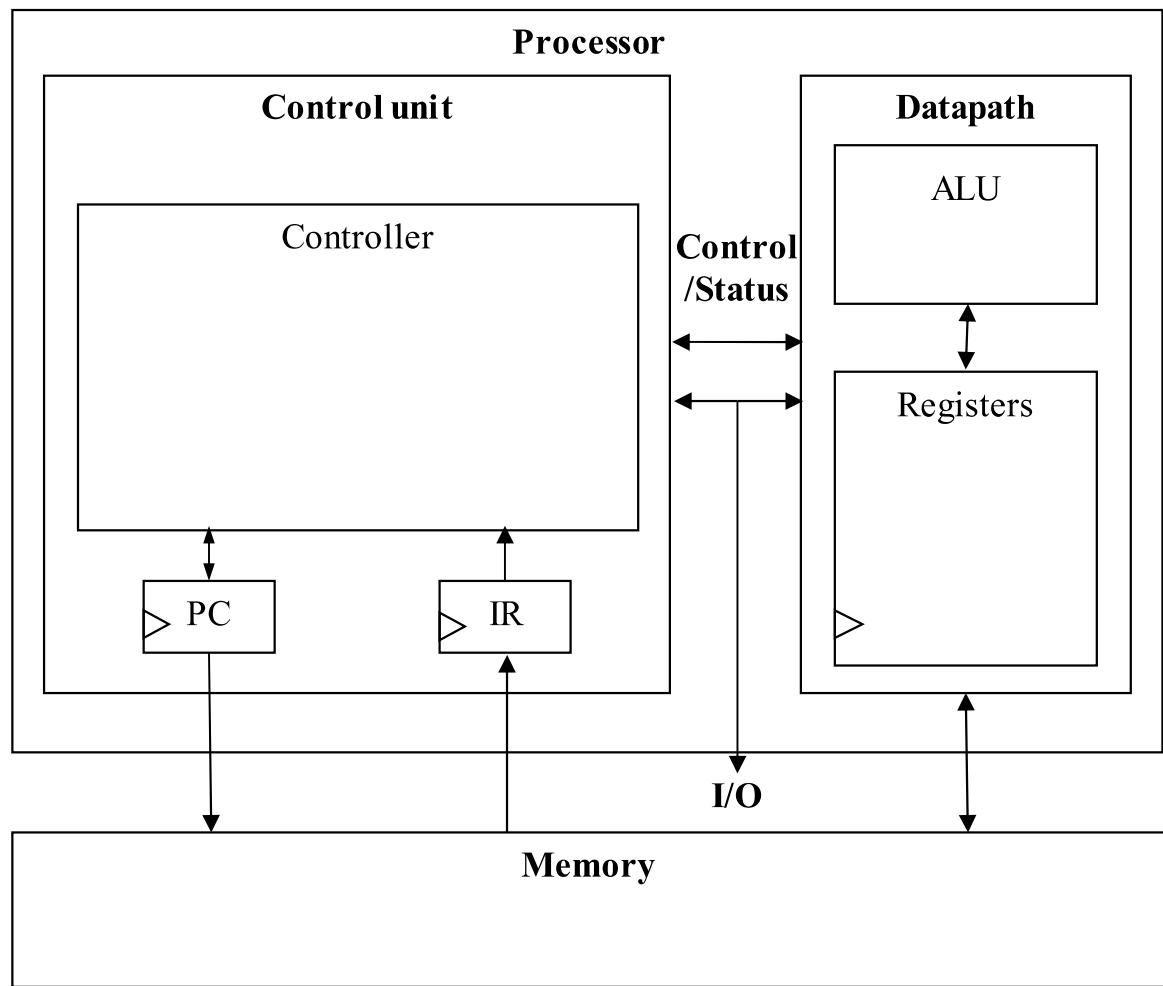


# Instruction Cycles



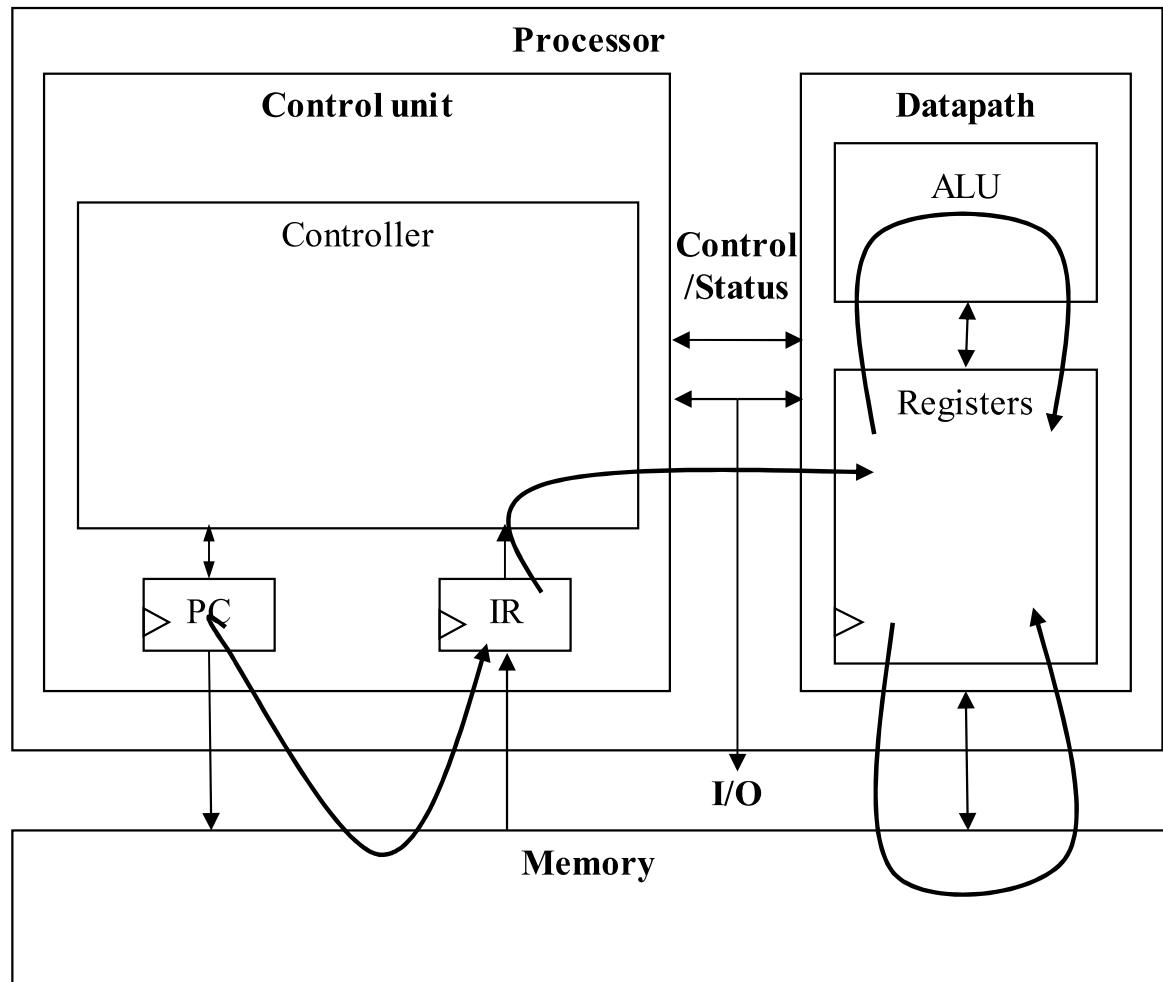
# Architectural Considerations

- $N$ -bit processor
  - $N$ -bit ALU, registers, buses, memory data interface
  - Embedded: 8-bit, 16-bit, 32-bit common
  - Desktop/servers: 32-bit, even 64
- PC size determines address space

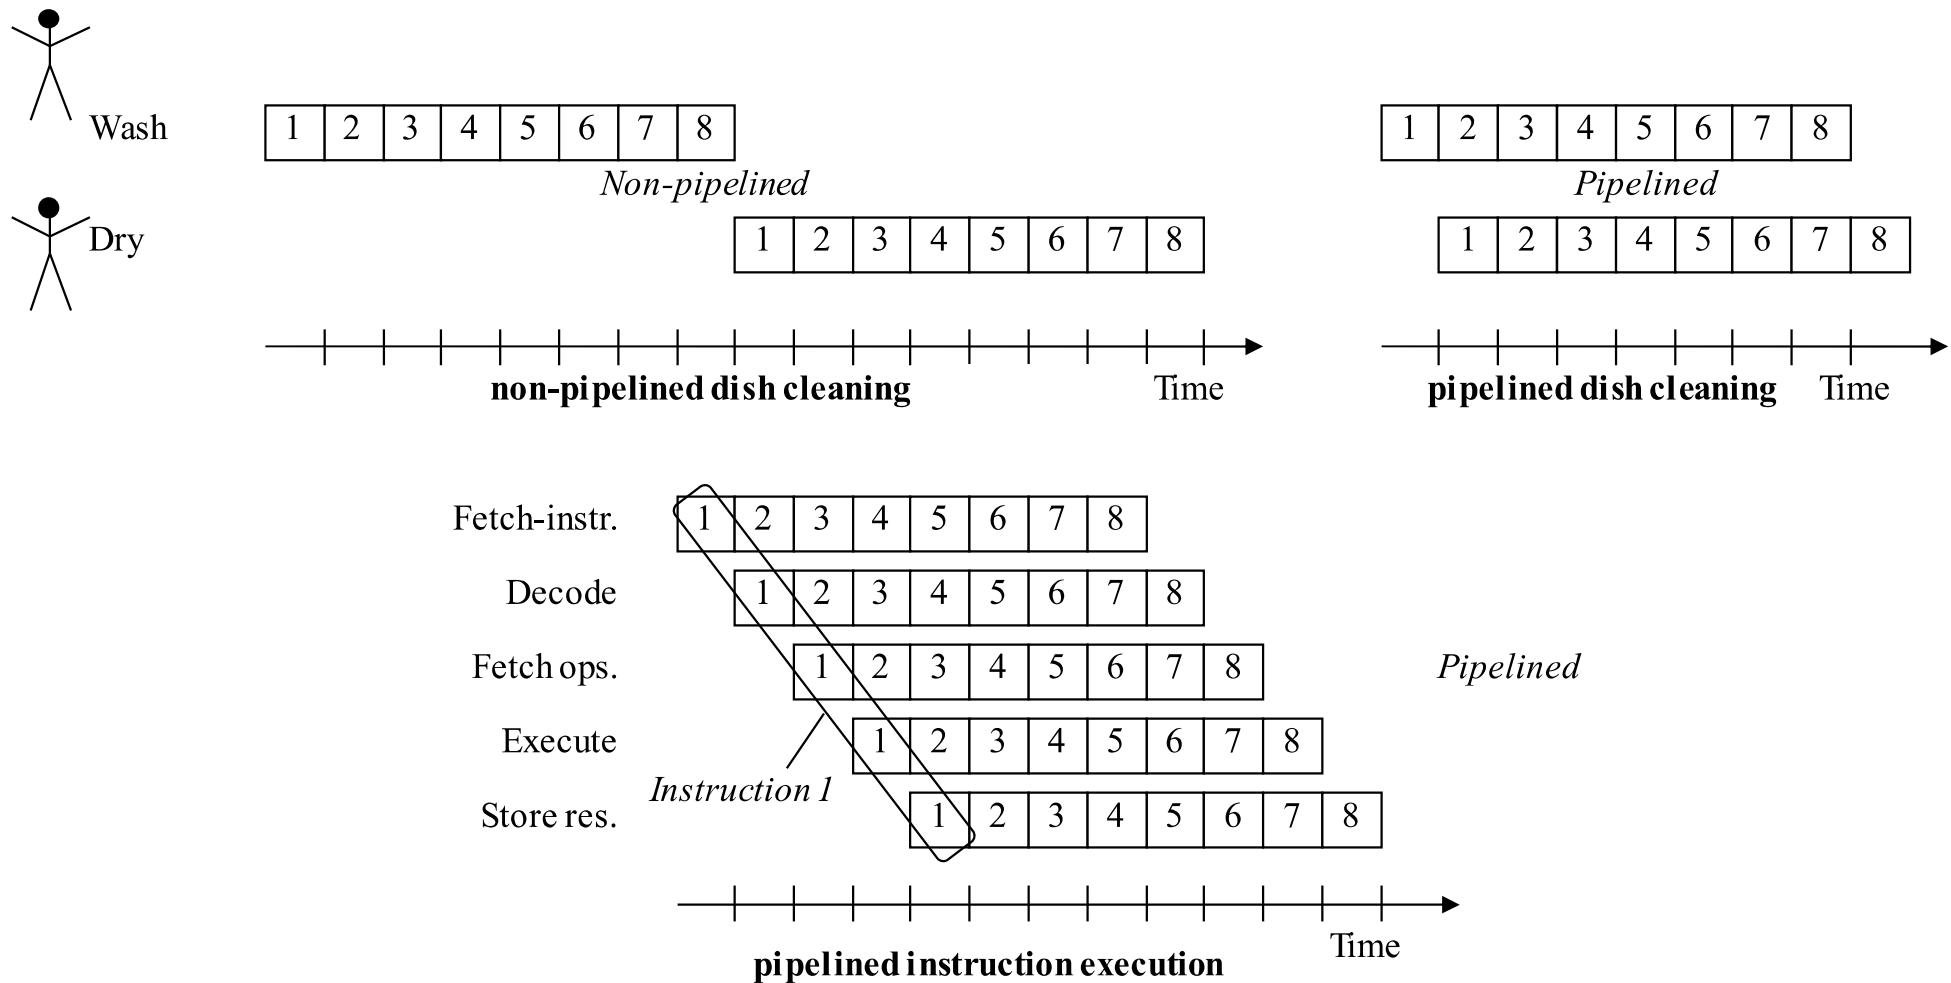


# Architectural Considerations

- Clock frequency
  - Inverse of clock period
  - Must be longer than longest register to register delay in entire processor
  - Memory access is often the longest



# Pipelining: Increasing Instruction *Throughput*



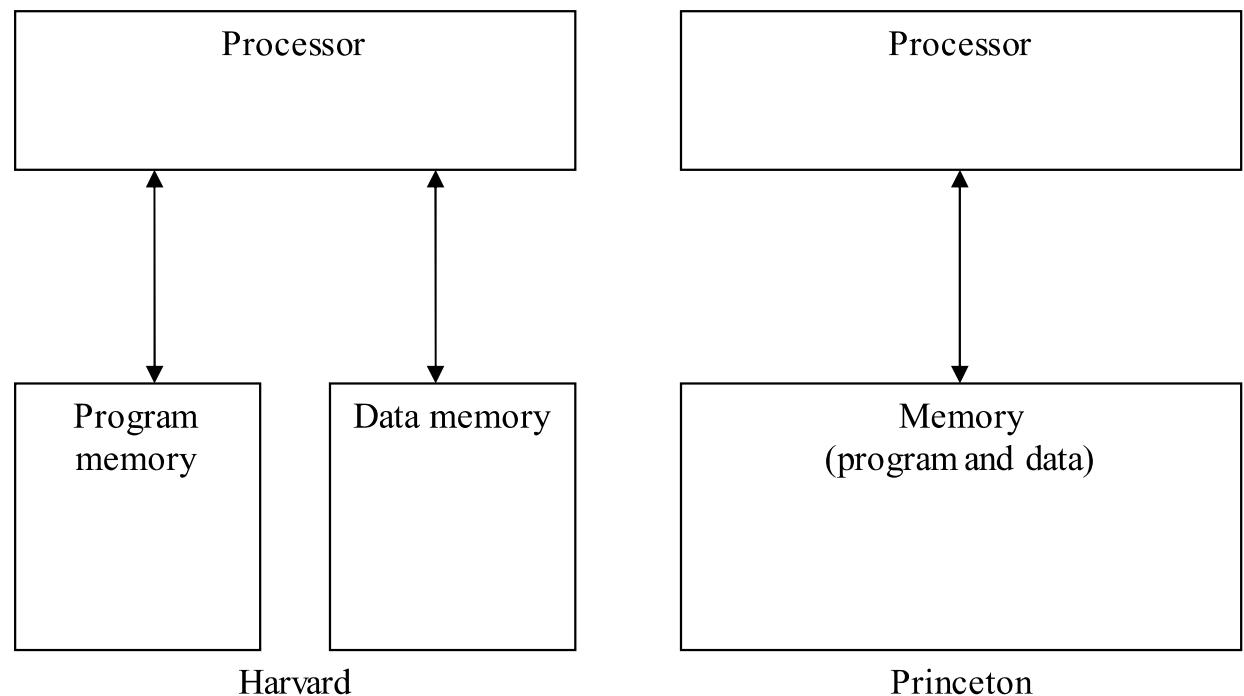
# Superscalar and VLIW Architectures

---

- Performance can be improved by:
  - Faster clock (but there's a limit)
  - Pipelining: slice up instruction into stages, overlap stages
  - *Multiple ALUs* to support more than one instruction stream
    - Superscalar
      - Scalar: non-vector operations
      - Fetches instructions in batches, executes as many as possible
        - May require extensive hardware to detect independent instructions
      - VLIW: each word in memory has multiple independent instructions
        - Relies on the compiler to detect and schedule instructions
        - Currently growing in popularity

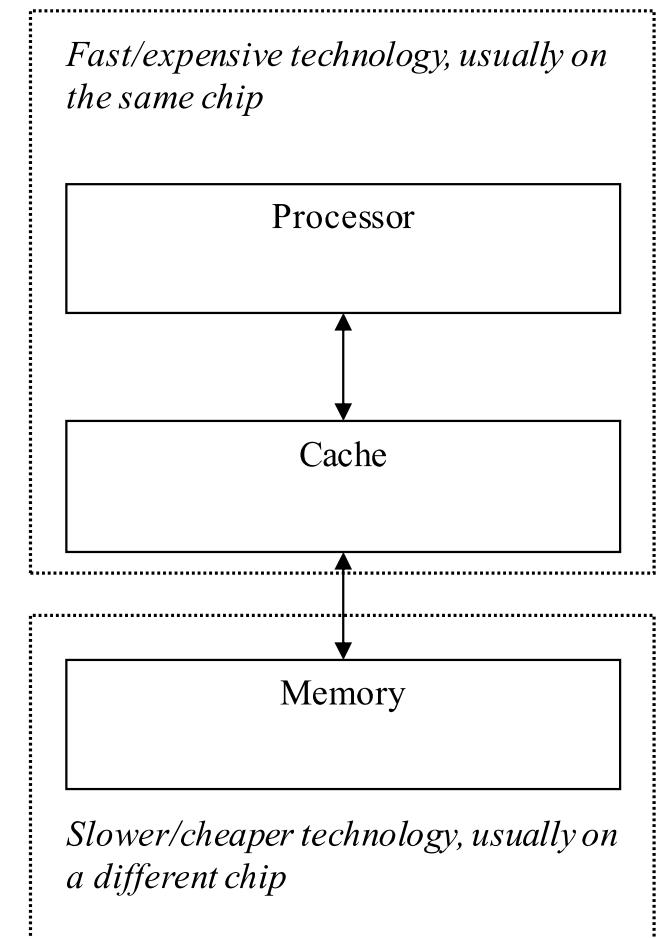
# Two Memory Architectures

- Princeton
  - Fewer memory wires
- Harvard
  - Simultaneous program and data memory access



# Cache Memory

- Memory access may be slow
- Cache is small but fast memory close to processor
  - Holds copy of part of memory
  - Hits and misses



# Programmer's View

---

- Programmer doesn't need detailed understanding of architecture
  - Instead, needs to know what instructions can be executed
- Two levels of instructions:
  - Assembly level
  - Structured languages (C, C++, Java, etc.)
- Most development today done using structured languages
  - But, some assembly level programming may still be necessary
  - Drivers: portion of program that communicates with and/or controls (drives) another device
    - Often have detailed timing considerations, extensive bit manipulation
    - Assembly level may be best for these

# Assembly-Level Instructions

Instruction 1	opcode	operand1	operand2
Instruction 2	opcode	operand1	operand2
Instruction 3	opcode	operand1	operand2
Instruction 4	opcode	operand1	operand2
...			

- Instruction Set
  - Defines the legal set of instructions for that processor
    - Data transfer: memory/register, register/register, I/O, etc.
    - Arithmetic/logical: move register through ALU and back
    - Branches: determine next PC value when not just PC+1

# A Simple (Trivial) Instruction Set

Assembly instruct.	First byte	Second byte	Operation				
MOV Rn, direct	<table border="1"><tr><td>0000</td><td>Rn</td></tr></table>	0000	Rn	<table border="1"><tr><td>direct</td></tr></table>	direct	$Rn = M(\text{direct})$	
0000	Rn						
direct							
MOV direct, Rn	<table border="1"><tr><td>0001</td><td>Rn</td></tr></table>	0001	Rn	<table border="1"><tr><td>direct</td></tr></table>	direct	$M(\text{direct}) = Rn$	
0001	Rn						
direct							
MOV @Rn, Rm	<table border="1"><tr><td>0010</td><td>Rn</td></tr></table>	0010	Rn	<table border="1"><tr><td>Rm</td><td></td></tr></table>	Rm		$M(Rn) = Rm$
0010	Rn						
Rm							
MOV Rn, #immed.	<table border="1"><tr><td>0011</td><td>Rn</td></tr></table>	0011	Rn	<table border="1"><tr><td>immediate</td></tr></table>	immediate	$Rn = \text{immediate}$	
0011	Rn						
immediate							
ADD Rn, Rm	<table border="1"><tr><td>0100</td><td>Rn</td></tr></table>	0100	Rn	<table border="1"><tr><td>Rm</td><td></td></tr></table>	Rm		$Rn = Rn + Rm$
0100	Rn						
Rm							
SUB Rn, Rm	<table border="1"><tr><td>0101</td><td>Rn</td></tr></table>	0101	Rn	<table border="1"><tr><td>Rm</td><td></td></tr></table>	Rm		$Rn = Rn - Rm$
0101	Rn						
Rm							
JZ Rn, relative	<table border="1"><tr><td>0110</td><td>Rn</td></tr></table>	0110	Rn	<table border="1"><tr><td>relative</td></tr></table>	relative	$PC = PC + \text{relative}$ (only if Rn is 0)	
0110	Rn						
relative							

# Addressing Modes

Addressing mode	Operand field	Register-file contents	Memory contents
Immediate	Data		
Register-direct	Register address	Data	
Register indirect	Register address	Memory address	Data
Direct	Memory address	.....	Data
Indirect	Memory address		Memory address
			Data

# Sample Programs

C program	Equivalent assembly program
	0 MOV R0,#0; // total = 0
	1 MOV R1,#10; // i = 10
	2 MOV R2,#1; // constant 1
	3 MOV R3,#0; // constant 0
int total=0;	Loop: JZ R1, Next; // Done if i=0
for (int i=10; i!=0; i--)	5 ADD R0, R1; // total += i
total += i;	6 SUB R1, R2; // i--
// next instructions...	7 JZ R3, Loop; // Jump always
	Next: // next instructions...

- Try some others
  - Handshake: Wait until the value of M[254] is not 0, set M[255] to 1, wait until M[254] is 0, set M[255] to 0 (assume those locations are ports).
  - (Harder) Count the occurrences of zero in an array stored in memory locations 100 through 199.

# Programmer Considerations

---

- Program and data memory space
  - Embedded processors often very limited
    - e.g., 64 Kbytes program, 256 bytes of RAM (expandable)
- Registers: How many are there?
  - Only a direct concern for assembly-level programmers
- I/O
  - How communicate with external signals?
- Interrupts

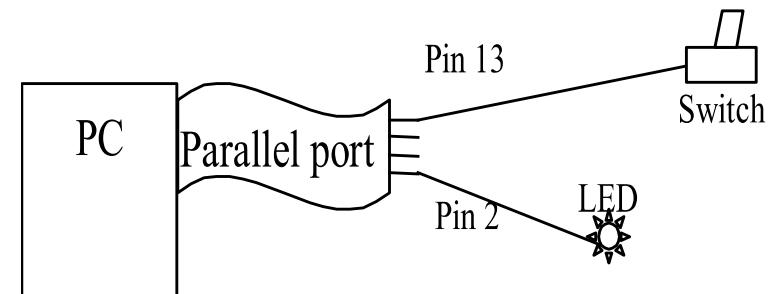
# Microprocessor Architecture Overview

---

- If you are using a particular microprocessor, now is a good time to review its architecture

# Example: parallel port driver

LPT Connection Pin	I/O Direction	Register Address
1	Output	0 <sup>th</sup> bit of register #2
2-9	Output	0 <sup>th</sup> bit of register #2
10,11,12,13,15	Input	6,7,5,4,3 <sup>th</sup> bit of register #1
14,16,17	Output	1,2,3 <sup>th</sup> bit of register #2



- Using assembly language programming we can configure a PC parallel port to perform digital I/O
  - write and read to three special registers to accomplish this table provides list of parallel port connector pins and corresponding register location
  - Example : parallel port monitors the input switch and turns the LED on/off accordingly

# Parallel Port Example

```
; This program consists of a sub-routine that reads
; the state of the input pin, determining the on/off state
; of our switch and asserts the output pin, turning the LED
; on/off accordingly
```

.386

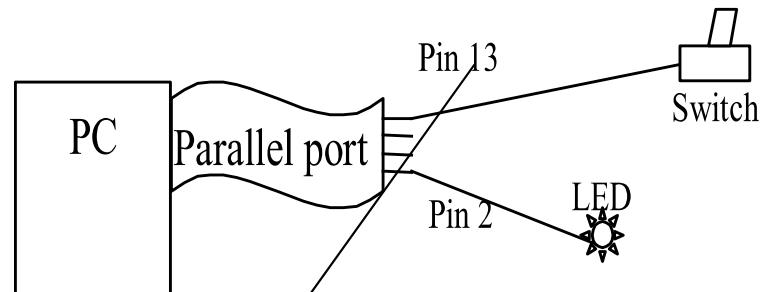
```
CheckPort proc
    push ax      ; save the content
    push dx      ; save the content
    mov dx, 3BCh + 1 ; base + 1 for register #1
    in al, dx    ; read register #1
    and al, 10h   ; mask out all but bit # 4
    cmp al, 0     ; is it 0?
    jne SwitchOn ; if not, we need to turn the LED on

SwitchOff:
    mov dx, 3BCh + 0 ; base + 0 for register #0
    in al, dx    ; read the current state of the port
    and al, f7h   ; clear first bit (masking)
    out dx, al   ; write it out to the port
    jmp Done     ; we are done

SwitchOn:
    mov dx, 3BCh + 0 ; base + 0 for register #0
    in al, dx    ; read the current state of the port
    or al, 01h    ; set first bit (masking)
    out dx, al   ; write it out to the port

Done:  pop dx      ; restore the content
    pop ax      ; restore the content
CheckPort endp
```

```
extern "C" CheckPort(void);           // defined in
                                         // assembly
void main(void) {
    while( 1 ) {
        CheckPort();
    }
}
```



LPT Connection Pin	I/O Direction	Register Address
1	Output	0 <sup>th</sup> bit of register #2
2-9	Output	0 <sup>th</sup> bit of register #2
10,11,12,13,15	Input	6,7,5,4,3 <sup>rd</sup> bit of register #1
14,16,17	Output	1,2,3 <sup>rd</sup> bit of register #2

# Operating System

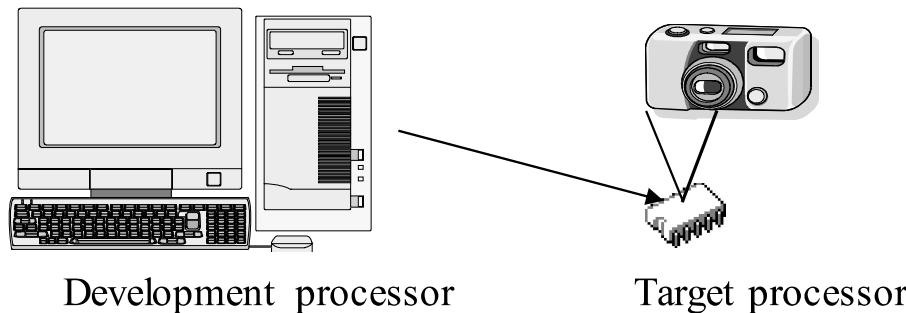
- Optional software layer providing low-level services to a program (application).
  - File management, disk access
  - Keyboard/display interfacing
  - Scheduling multiple programs for execution
    - Or even just multiple threads from one program
  - Program makes system calls to the OS

```
DB file_name "out.txt" -- store file name  
  
MOV R0, 1324           -- system call "open" id  
MOV R1, file_name      -- address of file-name  
INT 34                 -- cause a system call  
JZ  R0, L1              -- if zero -> error  
  
    . . . read the file  
JMP L2                  -- bypass error cond.  
L1:  
    . . . handle the error  
L2:
```

# Development Environment

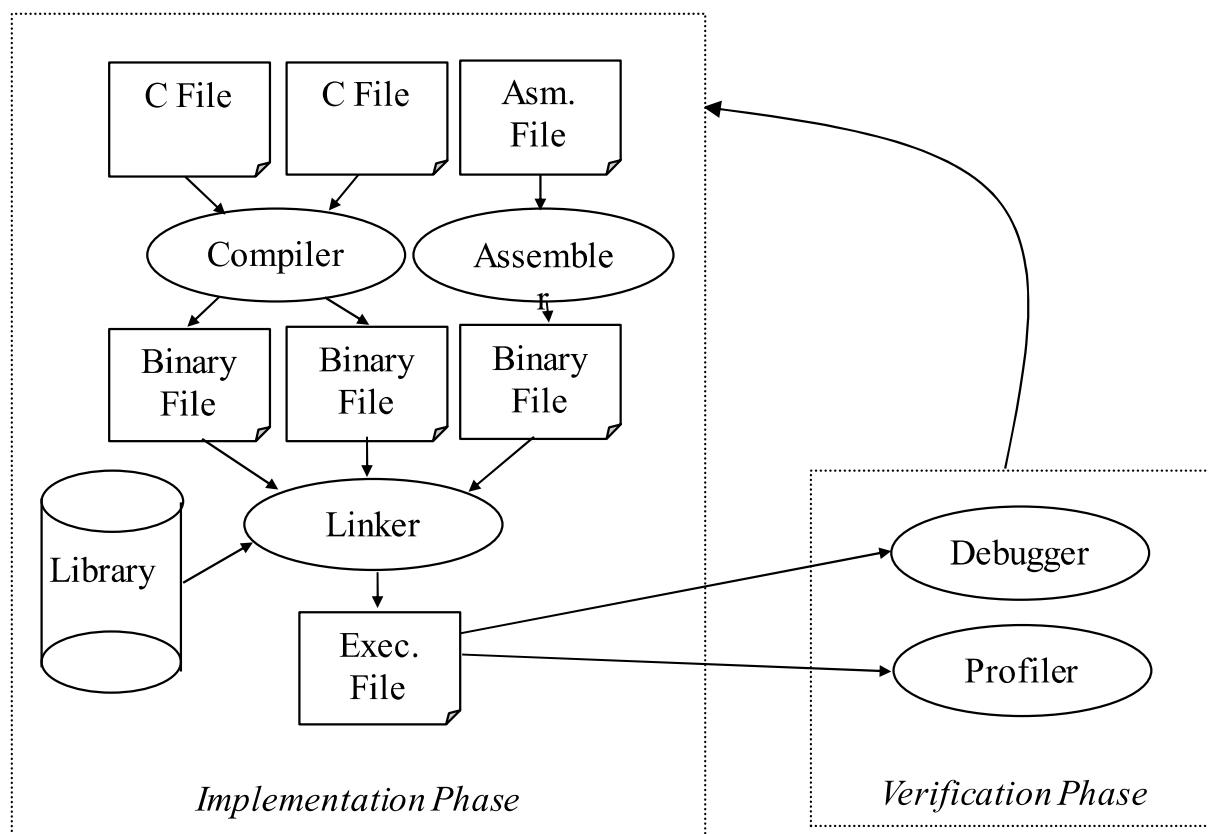
---

- Development processor
  - The processor on which we write and debug our programs
    - Usually a PC
- *Target processor*
  - The processor that the program will run on in our embedded system
    - Often different from the development processor



# Software Development Process

- Compilers
  - Cross compiler
    - Runs on one processor, but generates code for another
- Assemblers
- Linkers
- Debuggers
- Profilers



# Running a Program

---

- If development processor is different than target, how can we run our compiled code? Two options:
  - Download to target processor
  - Simulate
- Simulation
  - One method: Hardware description language
    - But slow, not always available
  - Another method: *Instruction set simulator (ISS)*
    - Runs on development processor, but executes instructions of target processor

# Instruction Set Simulator For A Simple Processor

```
#include <stdio.h>
typedef struct {
    unsigned char first_byte, second_byte;
} instruction;

instruction program[1024]; //instruction memory
unsigned char memory[256]; //data memory

void run_program(int num_bytes) {

    int pc = -1;
    unsigned char reg[16], fb, sb;

    while( ++pc < (num_bytes / 2) ) {
        fb = program[pc].first_byte;
        sb = program[pc].second_byte;
        switch( fb >> 4 ) {
            case 0: reg[fb & 0x0f] = memory[sb]; break;
            case 1: memory[sb] = reg[fb & 0x0f]; break;
            case 2: memory[reg[fb & 0x0f]] =
                reg[sb >> 4]; break;
            case 3: reg[fb & 0x0f] = sb; break;
            case 4: reg[fb & 0x0f] += reg[sb >> 4]; break;
            case 5: reg[fb & 0x0f] -= reg[sb >> 4]; break;
            case 6: pc += sb; break;
            default: return -1;
        }
    }
}
```

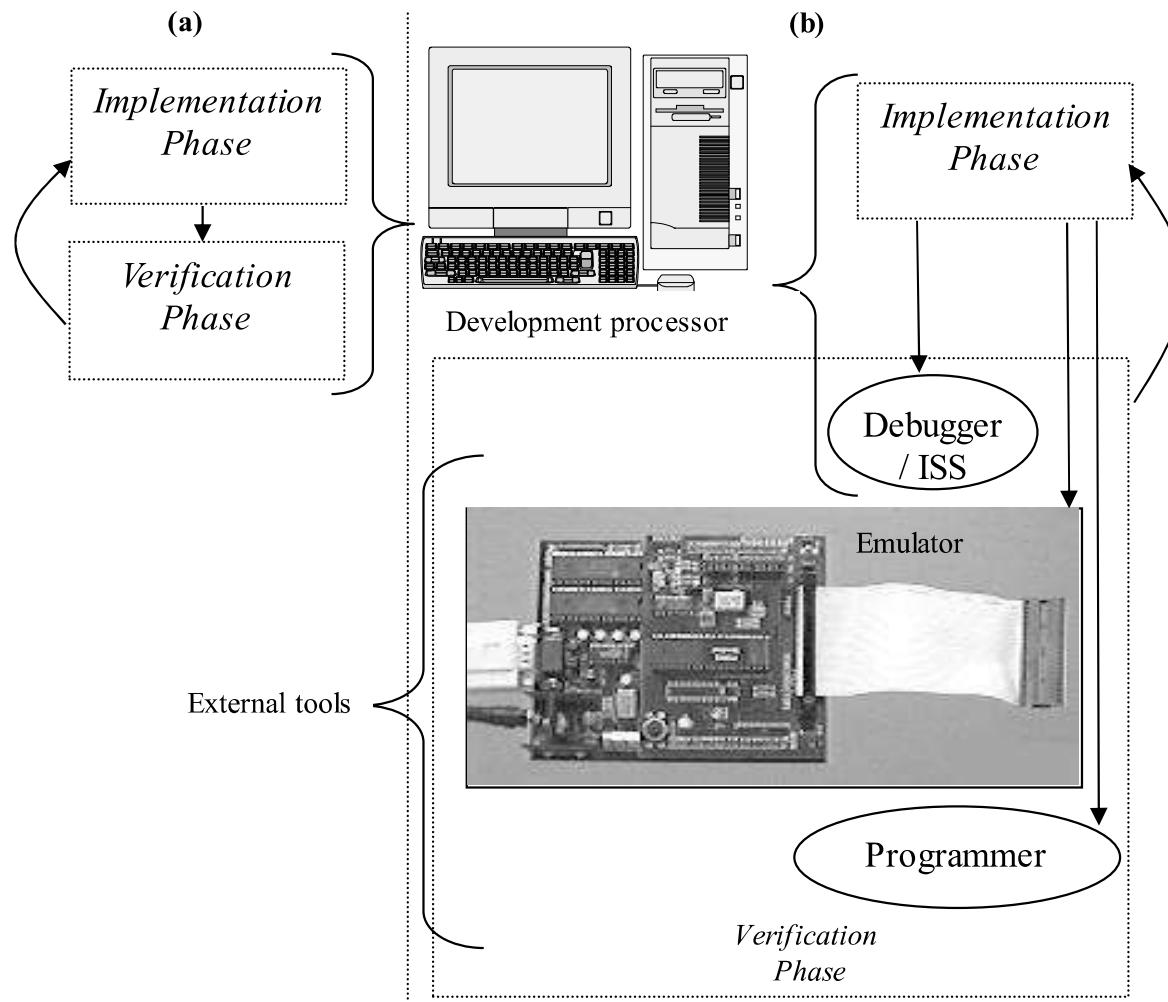
```
}
}

int main(int argc, char *argv[]) {

    FILE* ifs;

    If( argc != 2 ||
        (ifs = fopen(argv[1], "rb")) == NULL ) {
        return -1;
    }
    if (run_program(fread(program,
        sizeof(program) == 0) {
        print_memory_contents();
        return(0);
    }
    else return(-1);
}
```

# Testing and Debugging



- **ISS**
  - Gives us control over time – set breakpoints, look at register values, set values, step-by-step execution, ...
  - But, doesn't interact with real environment
- **Download to board**
  - Use device programmer
  - Runs in real environment, but not controllable
- **Compromise: emulator**
  - Runs in real environment, at speed or near
  - Supports some controllability from the PC

# Application-Specific Instruction-Set Processors (ASIPs)

---

- General-purpose processors
  - Sometimes too general to be effective in demanding application
    - e.g., video processing – requires huge video buffers and operations on large arrays of data, inefficient on a GPP
  - But single-purpose processor has high NRE, not programmable
- ASIPs – targeted to a particular domain
  - Contain architectural features specific to that domain
    - e.g., embedded control, digital signal processing, video processing, network processing, telecommunications, etc.
  - Still programmable

# A Common ASIP: Microcontroller

---

- For embedded control applications
  - Reading sensors, setting actuators
  - Mostly dealing with events (bits): data is present, but not in huge amounts
  - e.g., VCR, disk drive, digital camera (assuming SPP for image compression), washing machine, microwave oven
- Microcontroller features
  - On-chip peripherals
    - Timers, analog-digital converters, serial communication, etc.
    - Tightly integrated for programmer, typically part of register space
  - On-chip program and data memory
  - Direct programmer access to many of the chip's pins
  - Specialized instructions for bit-manipulation and other low-level operations

# Another Common ASIP: Digital Signal Processors (DSP)

---

- For signal processing applications
  - Large amounts of digitized data, often streaming
  - Data transformations must be applied fast
  - e.g., cell-phone voice filter, digital TV, music synthesizer
- DSP features
  - Several instruction execution units
  - Multiple-accumulate single-cycle instruction, other instrs.
  - Efficient vector operations – e.g., add two arrays
    - Vector ALUs, loop buffers, etc.

# Trend: Even More Customized ASIPs

---

- In the past, microprocessors were acquired as chips
- Today, we increasingly acquire a processor as Intellectual Property (IP)
  - e.g., synthesizable VHDL model
- Opportunity to add a custom datapath hardware and a few custom instructions, or delete a few instructions
  - Can have significant performance, power and size impacts
  - Problem: need compiler/debugger for customized ASIP
    - Remember, most development uses structured languages
    - One solution: automatic compiler/debugger generation
      - e.g., [www.tensilica.com](http://www.tensilica.com)
    - Another solution: retargetable compilers
      - e.g., [www.improvsys.com](http://www.improvsys.com) (customized VLIW architectures)

# Selecting a Microprocessor

---

- Issues
  - Technical: speed, power, size, cost
  - Other: development environment, prior expertise, licensing, etc.
- Speed: how evaluate a processor's speed?
  - Clock speed – but instructions per cycle may differ
  - Instructions per second – but work per instr. may differ
  - Dhrystone: Synthetic benchmark, developed in 1984. Dhrystones/sec.
    - MIPS: 1 MIPS = 1757 Dhrystones per second (based on Digital's VAX 11/780). A.k.a. Dhrystone MIPS. Commonly used today.
      - So, 750 MIPS =  $750 \times 1757 = 1,317,750$  Dhrystones per second
  - SPEC: set of more realistic benchmarks, but oriented to desktops
  - EEMBC – EDN Embedded Benchmark Consortium, [www.eembc.org](http://www.eembc.org)
    - Suites of benchmarks: automotive, consumer electronics, networking, office automation, telecommunications

# General Purpose Processors

Processor	Clock speed	Periph.	Bus Width	MIPS	Power	Trans.	Price
General Purpose Processors							
Intel PIII	1GHz	2x16 K L1, 256K L2, MMX	32	~900	97W	~7M	\$900
IBM PowerPC 750X	550 MHz	2x32 K L1, 256K L2	32/64	~1300	5W	~7M	\$900
MIPS R5000	250 MHz	2x32 K 2 way set assoc.	32/64	NA	NA	3.6M	NA
StrongARM SA-110	233 MHz	None	32	268	1W	2.1M	NA
Microcontroller							
Intel 8051	12 MHz	4K ROM, 128 RAM, 32 I/O, Timer, UART	8	~1	~0.2W	~10K	\$7
Motorola 68HC811	3 MHz	4K ROM, 192 RAM, 32 I/O, Timer, WDT, SPI	8	~.5	~0.1W	~10K	\$5
Digital Signal Processors							
TI C5416	160 MHz	128K, SRAM, 3 T1 Ports, DMA, 13 ADC, 9 DAC	16/32	~600	NA	NA	\$34
Lucent DSP32C	80 MHz	16K Inst., 2K Data, Serial Ports, DMA	32	40	NA	NA	\$75

Sources: Intel, Motorola, MIPS, ARM, TI, and IBM Website/Datasheet; *Embedded Systems Programming*, Nov. 1998

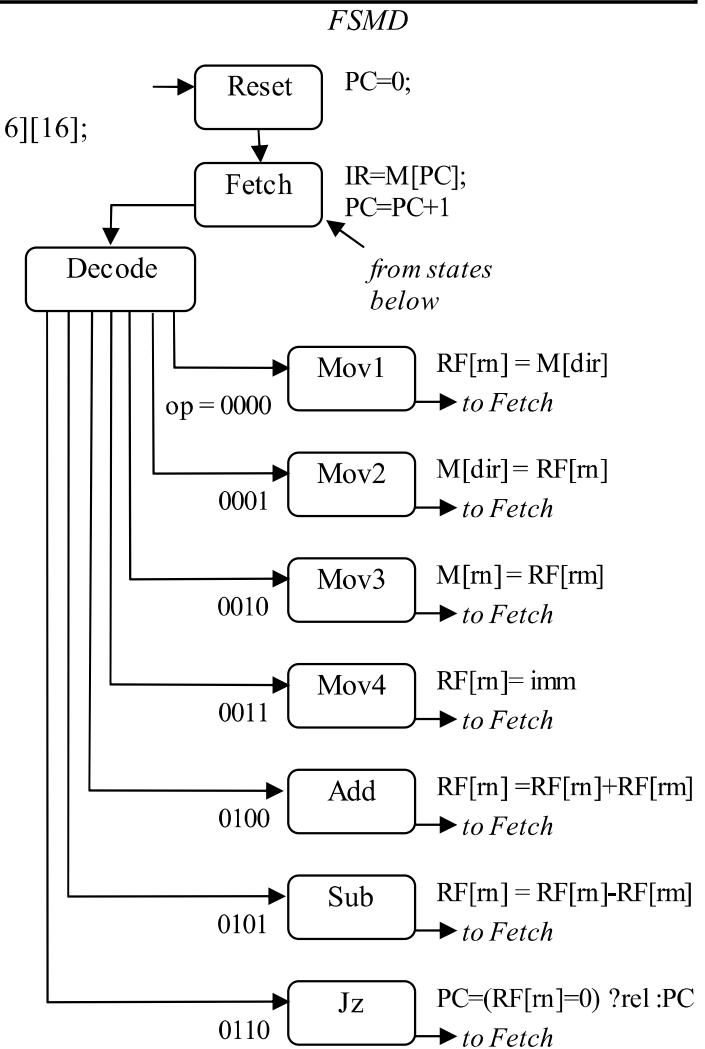
# Designing a General Purpose Processor

- Not something an embedded system designer normally would do
  - But instructive to see how simply we can build one top down
  - Remember that real processors aren't usually built this way
    - Much more optimized, much more bottom-up design

Aliases:	
op	IR[15..12]
rn	IR[11..8]
rm	IR[7..4]
dir	IR[7..0]
imm	IR[7..0]
rel	IR[7..0]

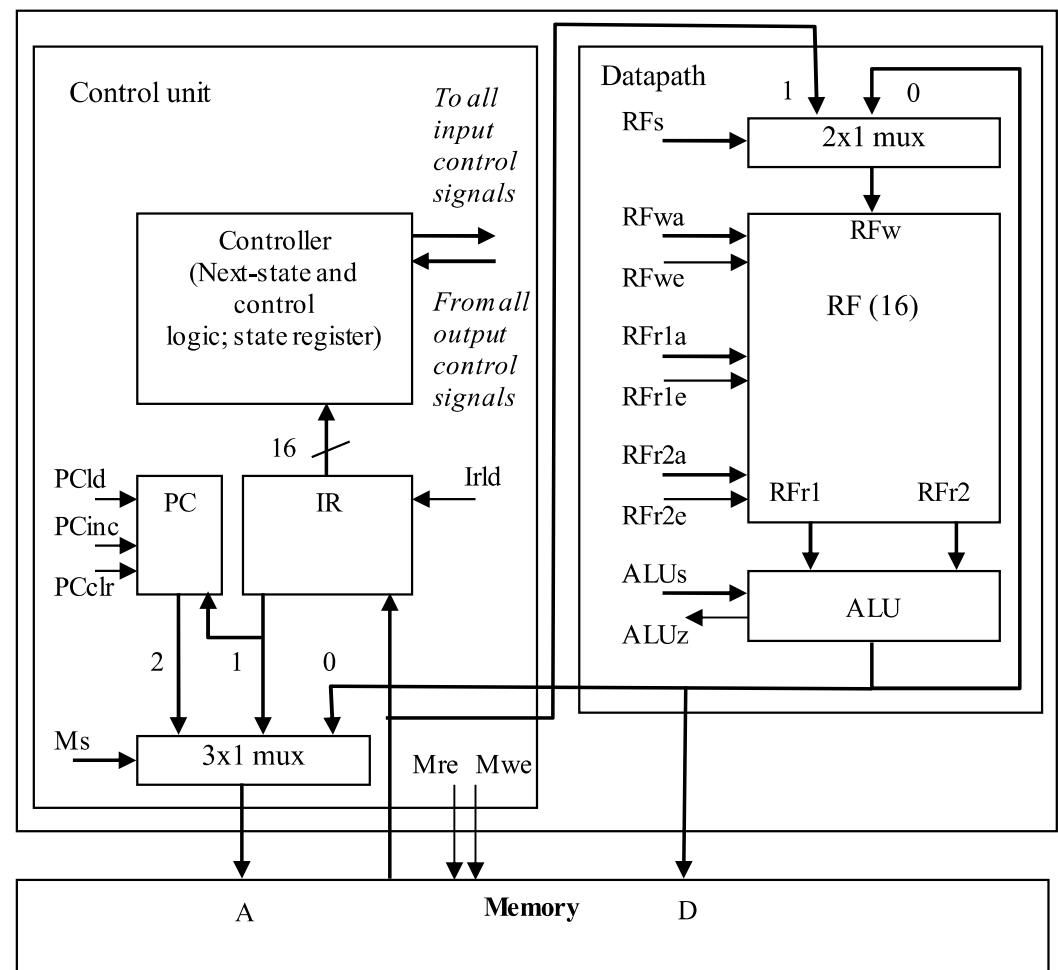
Declarations:

bit PC[16], IR[16];  
bit M[64k][16], RF[16][16];

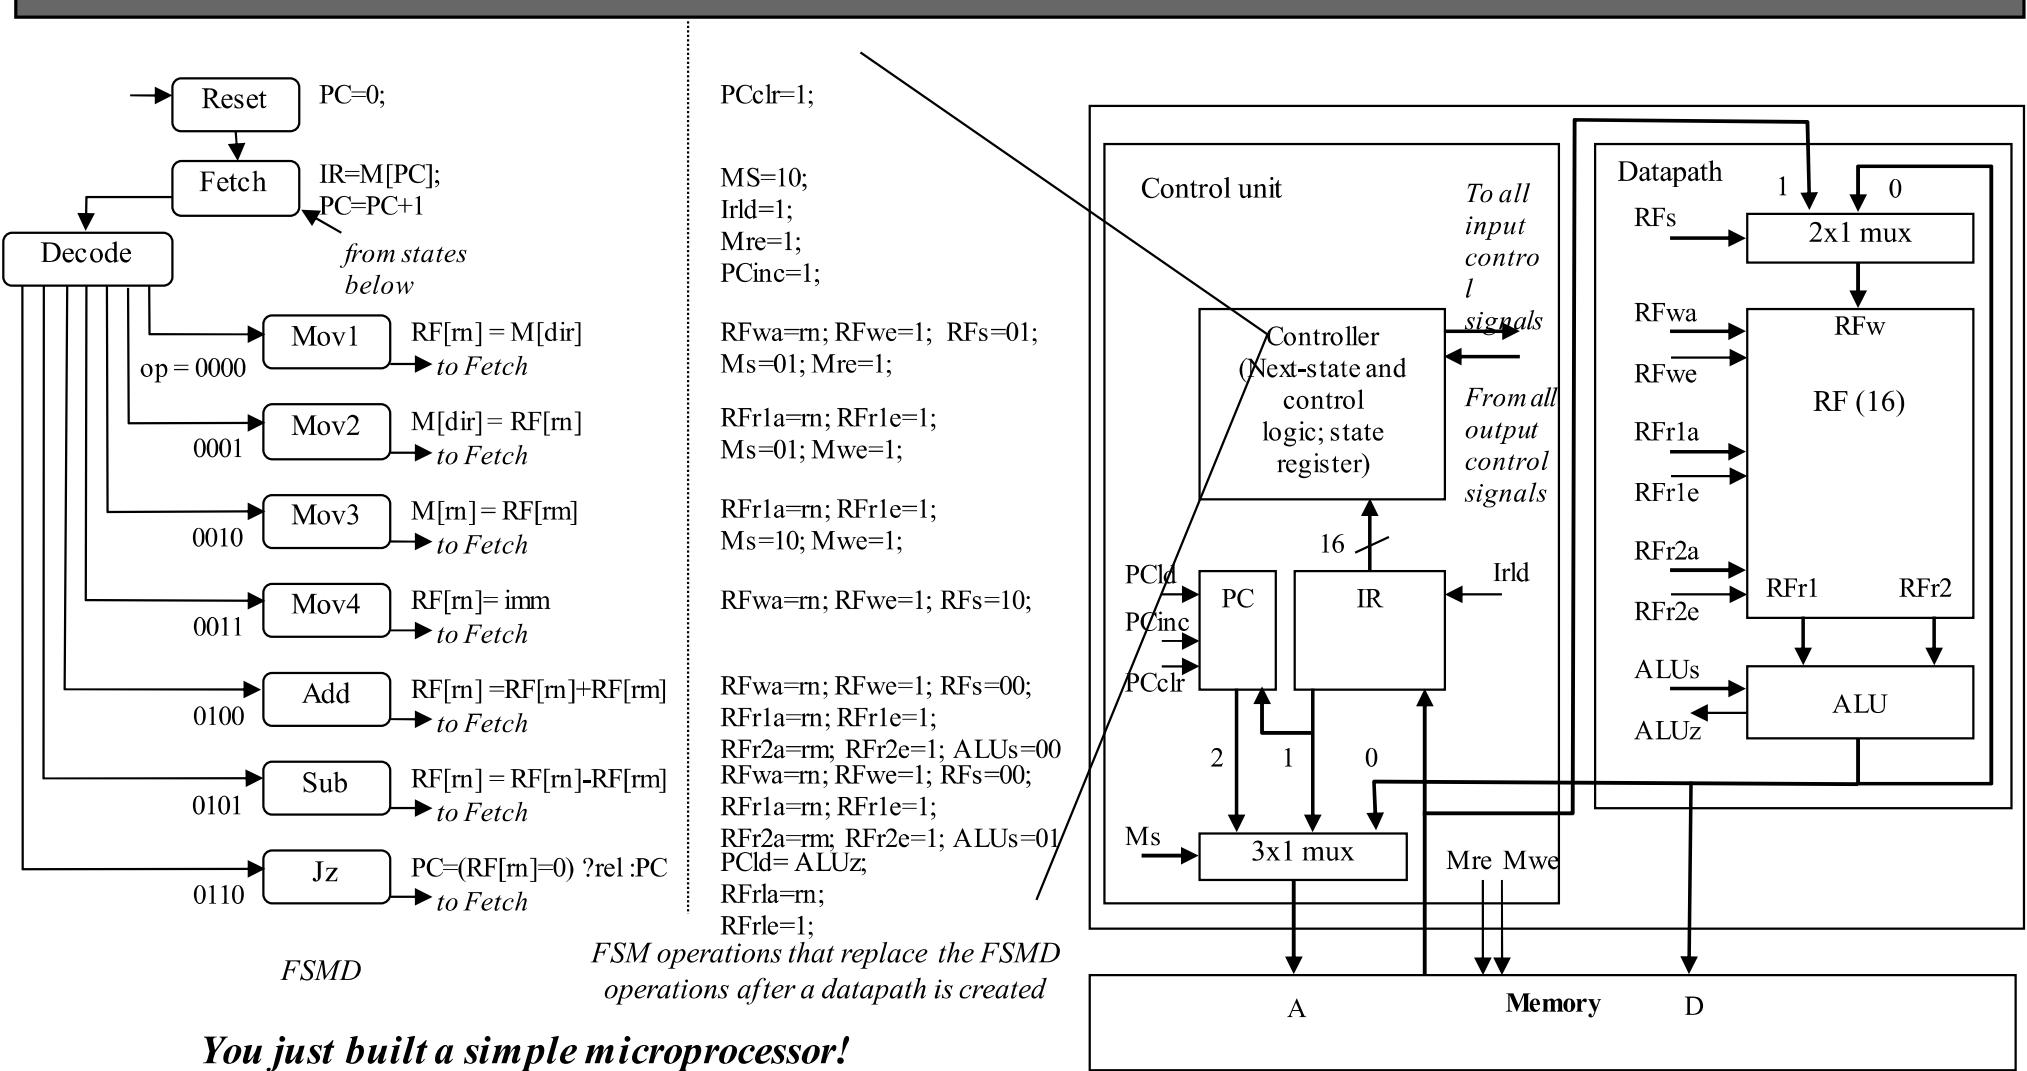


# Architecture of a Simple Microprocessor

- Storage devices for each declared variable
  - register file holds each of the variables
- Functional units to carry out the FSMD operations
  - One ALU carries out every required operation
- Connections added among the components' ports corresponding to the operations required by the FSM
- Unique identifiers created for every control signal



# A Simple Microprocessor



# Chapter Summary

---

- General-purpose processors
    - Good performance, low NRE, flexible
  - Controller, datapath, and memory
  - Structured languages prevail
    - But some assembly level programming still necessary
  - Many tools available
    - Including instruction-set simulators, and in-circuit emulators
  - ASIPs
    - Microcontrollers, DSPs, network processors, more customized ASIPs
  - Choosing among processors is an important step
  - Designing a general-purpose processor is conceptually the same as designing a single-purpose processor
-